

# SQLGraph

When ClickHouse marries graph processing

郑天祺 (Amos Bird)

中科院 计算所 网络数据实验室

# About Me

- Active ClickHouse Contributor
  - ~70 valid PRs
  - ~40 Stack Overflow Answers
- Open Source Enthusiast (Hacked 20+ Projects)
  - DB: Impala, Greenplum, Cockroach, Citus, Kudu
  - Misc: emacs, tmux, gdb, fish-shell, tdesktop ...
- SQLGraph Author




<https://github.com/amosbird>

# About My Lab

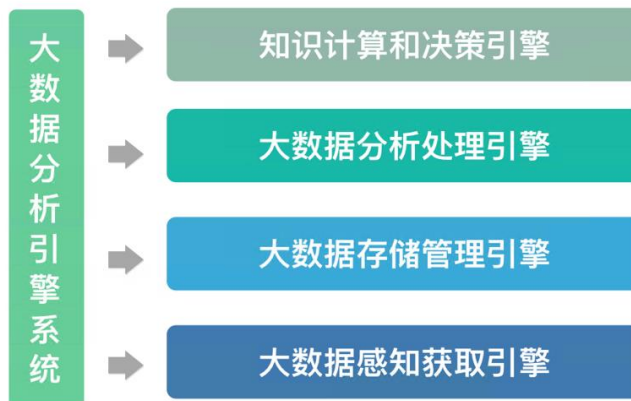
<http://www.ict.ac.cn/jgsz/kyxt/wlzdsys/>

- CCF Task Force on Big Data （大数据专家委员会）
- Organizing BDTC （大数据大会）
- Practitioner on Big Data



## 难点和目标:

- 1 解决分析深度浅的问题，提高数据的利用率
- 2 解决处理效率低的问题，提高计算的时效性
- 3 解决使用门槛高的问题，提高系统的易用性



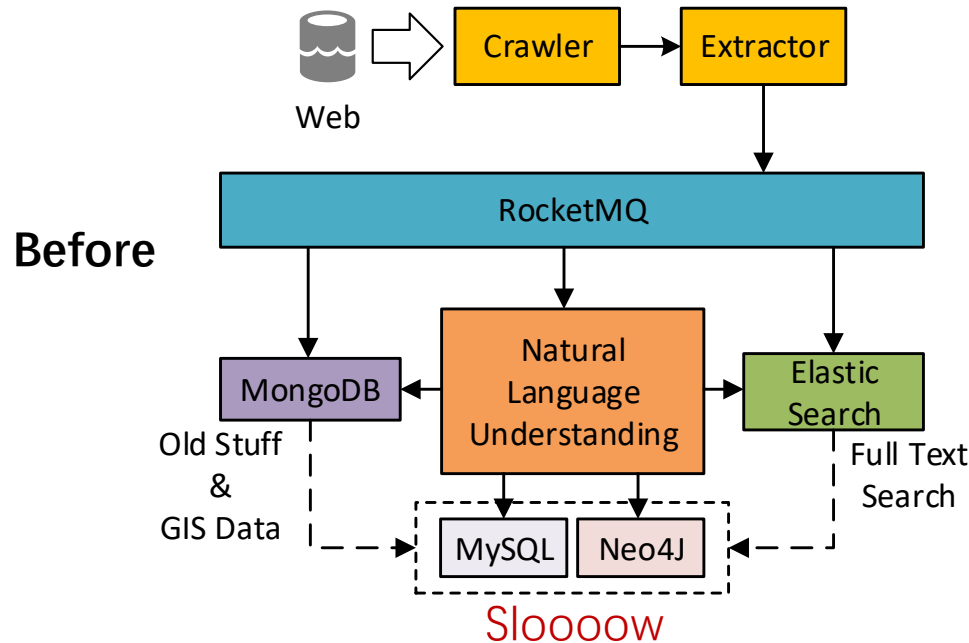
# About My Lab

<http://www.ict.ac.cn/jgsz/kyxt/wlzdsys/>

- CCF Task Force on Big Data （大数据专家委员会）
- Organizing BDTC （大数据大会）
- Practitioner on Big Data



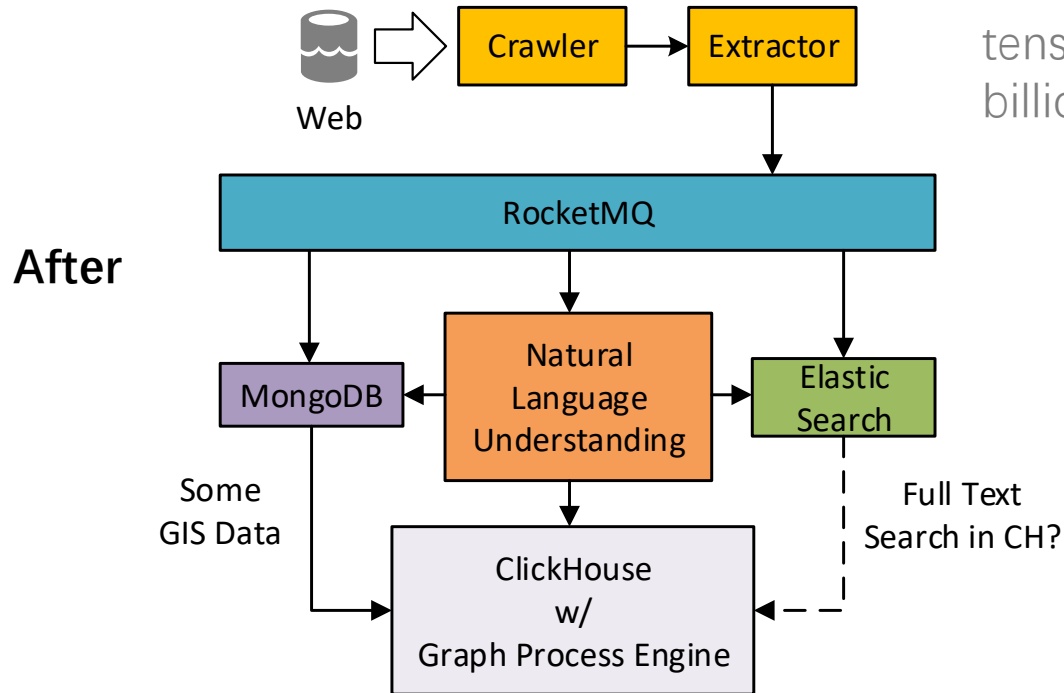
# How We Use ClickHouse



The NLU module outputs well structured, clean, (almost) immutable data.

*NLU : Natural Language Understanding*

# How We Use ClickHouse



Swift

Graph process engine is a specialized component to run analytical graph algorithms efficiently.


The NLU module outputs well structured, clean, (almost) immutable data.

*NLU : Natural Language Understanding*

# Why Graph Processing Engine

- Finding Important Nodes and Edges

- PageRank
- Personalized PageRank
- Shortest Path



- Community Detection

- Label Propagation
- Connection Components
- Louvain



- Recommendation

- Collaborative Filtering
- Random Walk
- Neural Embeddings



# Why Graph Processing Engine in RDBMS

- Hidden graph structures all over the place
  - Graph traversal via joins is both slow and inconvenient
  - Almost impossible to do iterative algorithms (recursion)
  - Cumbersome graph extraction process if acting as a graph data source
- Powerful in data management and processing
  - Full fledged processing (Instead of awk, perl ... )
  - Mature concepts of metadata handling
  - Widely adopted and ease of use interface



Mining what exists, aiding what emerges



# Why Graph Processing Engine in ClickHouse

- From Architect's Perspective
  - Fast data storage and fast data processing
  - Can handle many data sources
  - Rich user interface
- From Developer's Perspective
  - Nice building blocks for performance critical applications
  - Versatile SQL pipelines for mixed SQL/graph processing
  - Good code quality

Think OLAP but in the graph field.

# What to Expect from an OLAP Graph DB

- General Users
  - Easy to use (think as a graph)
  - Easy to understand (view as a graph)
  - Faster than Non-Graph databases or OLTP graph databases
- Graph Experts
  - Extremely fast analysis
  - Interactive graph algorithm designing
  - Ability to do low level optimizations

Be efficient to end users and computer hardware.

# System Design Goals

## **Usability**

1. Visualization
2. Graphs Being First-Class Citizens
3. SQL accessibility

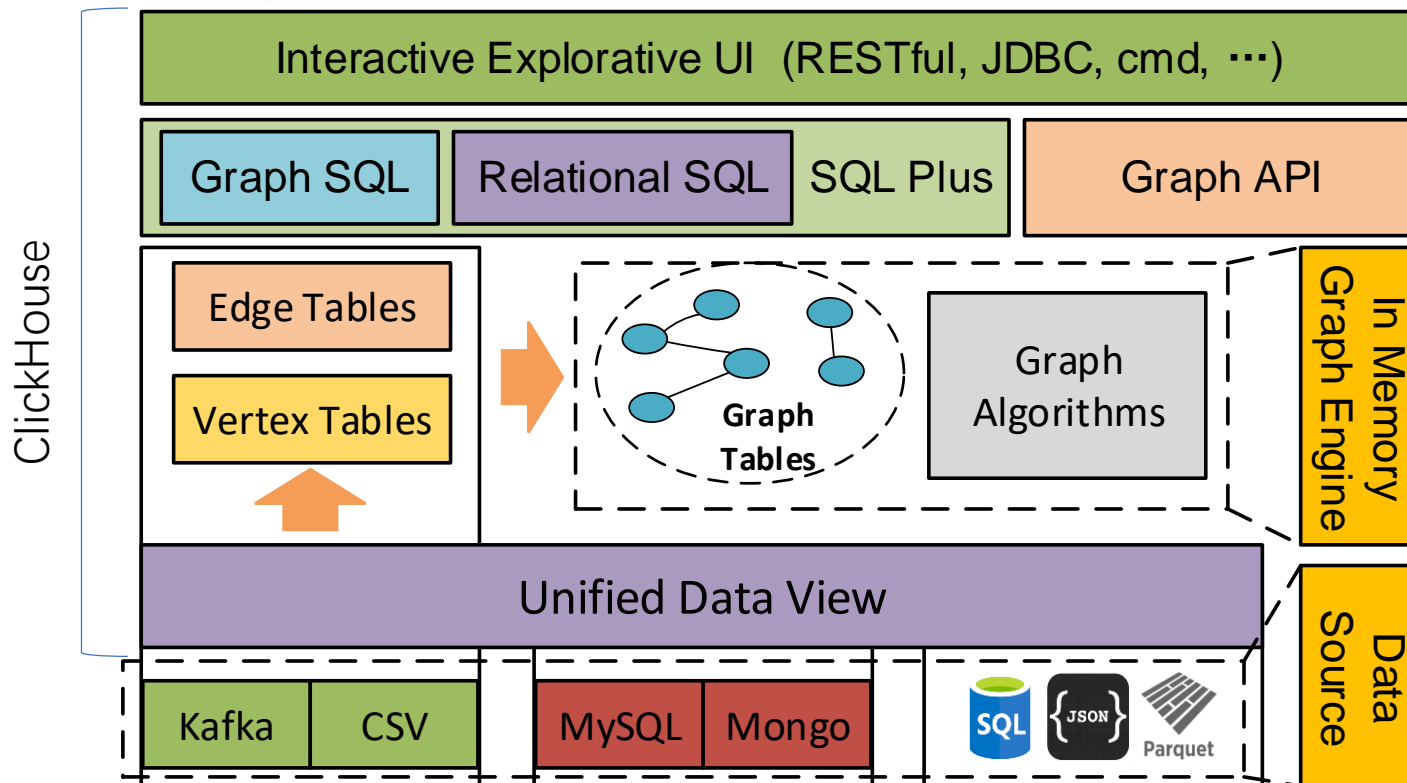

## **Versatile**

1. Graph Algorithm SDK
2. Graph Inspector and Rich Operators
3. Reusable Graph Modeling

## **Performance**

1. Outperform graph compute engines
2. Quick development iterations
3. Fast graph data digestion

# SQLGraph



# Some Results

## Calculate PageRank value per person

```
SELECT vp(v, 'name') AS name, pagerank
FROM pagerank(wz)
ORDER BY pagerank DESC LIMIT 5
```

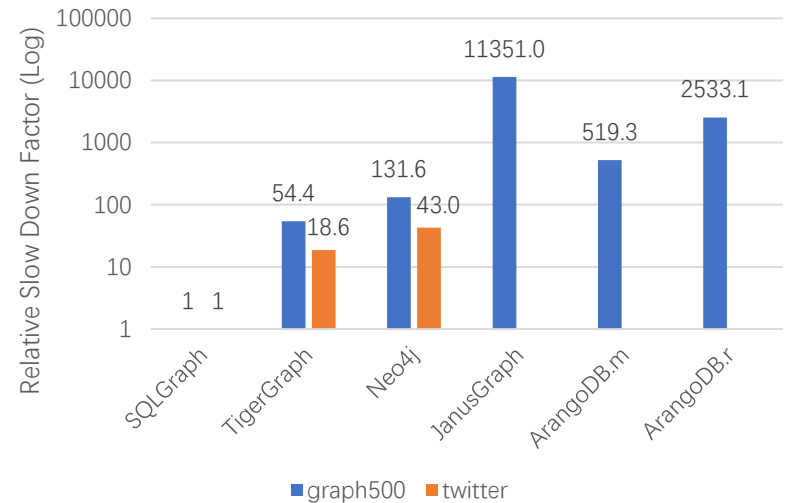
name	pagerank
li	0.06964007
zhao	0.063854694
qian	0.06365149
sun	0.06347877
shen	0.019985389

## Find a longest path which ends at 'shen'


```
SELECT pp(p, 'name', 'w') AS path
FROM edgePath(wz)
WHERE vp(p[-1].2, 'name') = 'shen'
ORDER BY length(path) DESC
LIMIT 1
```

path
zhou--[3]->wang--[2]->chu--[4]->shen

PageRank



Three-Hop Path Query



# Let's Dive in

examples and implementations

# Graph Query Interface

- Graph Algorithms are table functions
- Graphs are special tables that used as the first argument
- Collection of functions to retrieve graph info
- Example : Get the top 5 pagerank value from graph “wz”

```
SELECT
    vp(v, 'name'),      -- retrieve the vertex property “name”
    pagerank
FROM pagerank(          -- the graph algorithm
    wz,                 -- the graph table
    5,                  -- iterations
    0.85,               -- damping factor
    0.01                -- epsilon
)
ORDER BY pagerank DESC LIMIT 5
```

# Graph Query Output

- Three kinds of outputs : vertices, edges or graph info
- Vertices : PageRank, CC, BC, Radii, etc.
  - Has 'v : Vertices' column in the output list
  - Cover all vertices, no duplications
  - Can be inserted back into the graph as a global vertex property

```
SELECT
    id,
    groupArray(vp(v, 'name')) AS names
FROM cc(wz)
GROUP BY id
```

id	names
0	['zhao','qian','sun','li']
4	['zhou','wu','zheng','wang','feng','chen','chu','wei','shen']
14	['han','yang']
12	['jiang']



# Graph Query Output

- Three kinds of outputs : vertices, edges or graph info
- Edges : Hop, SSSP, CommonNeighbors, LinkCircle, etc.
  - Has 'e : Edges' column in the output list
  - Support function combinators: Out, In, All, Path
  - Can be used to derive subgraphs: create graph as select ...
  - Has 'Graph' format to support Non-Structural output (for visualization)

```
SELECT pp(p, 'name', 'w')
FROM hopInPath(wz,
(
  SELECT v
  FROM vertex(wz)
  WHERE name = 'zhao'
))
```

```
pp(p, 'name', 'w')
zhao<-[1]--li<-[2]--sun<-[3]--qian<-[2]--zhao
zhao<-[1]--li<-[4]--zhou
```

# Graph Query Output

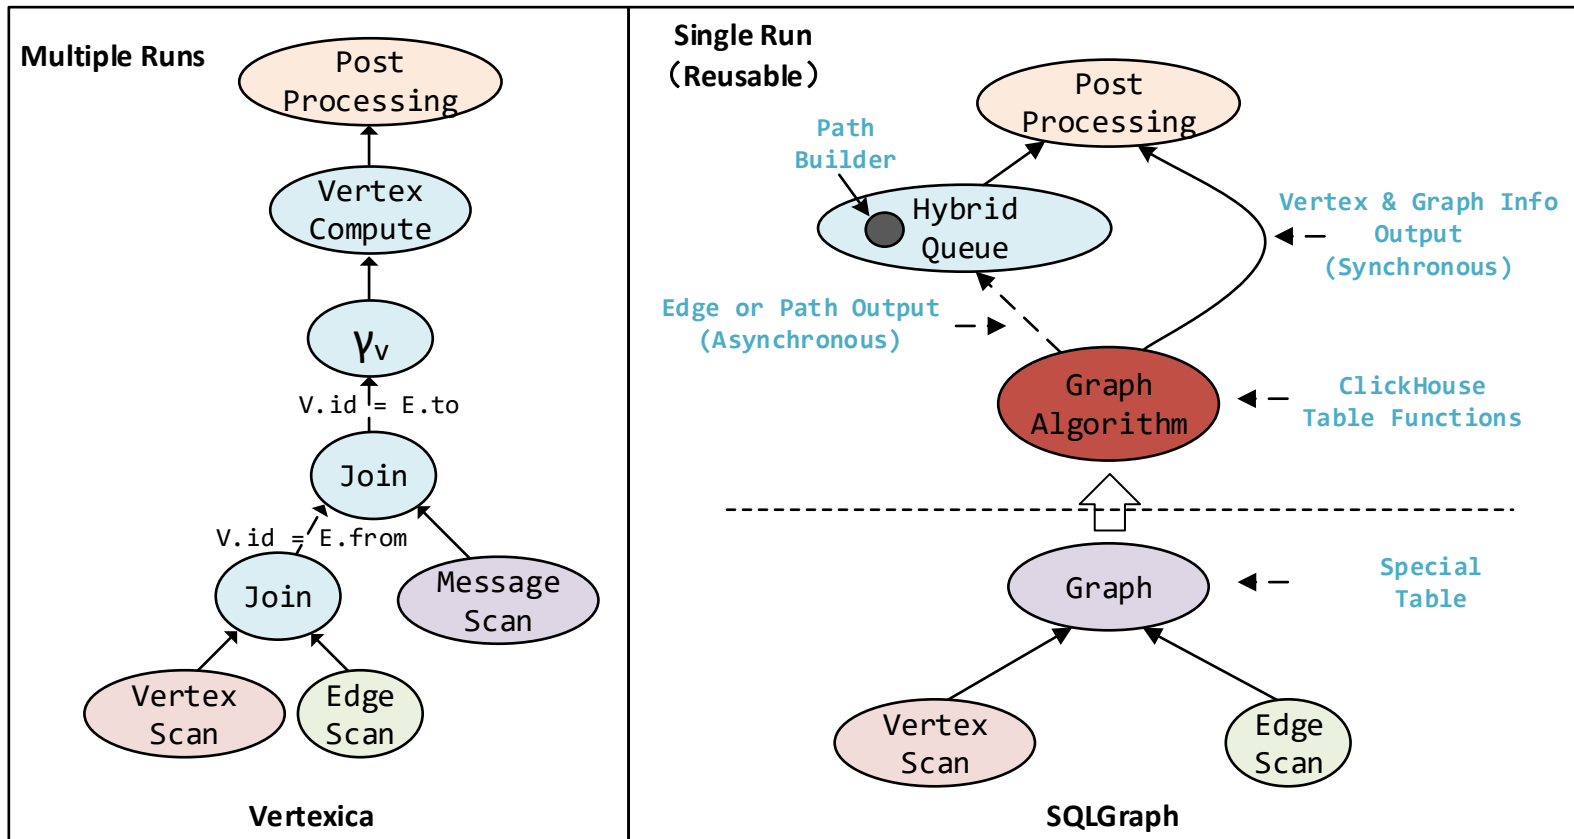
- Three kinds of outputs : vertices, edges or graph info
- Graph info : TriangleCount, MaxClique, ClusteringCoef, etc.
  - Results are graph characters

```
SELECT *  
FROM triangle(wz)
```

count
2

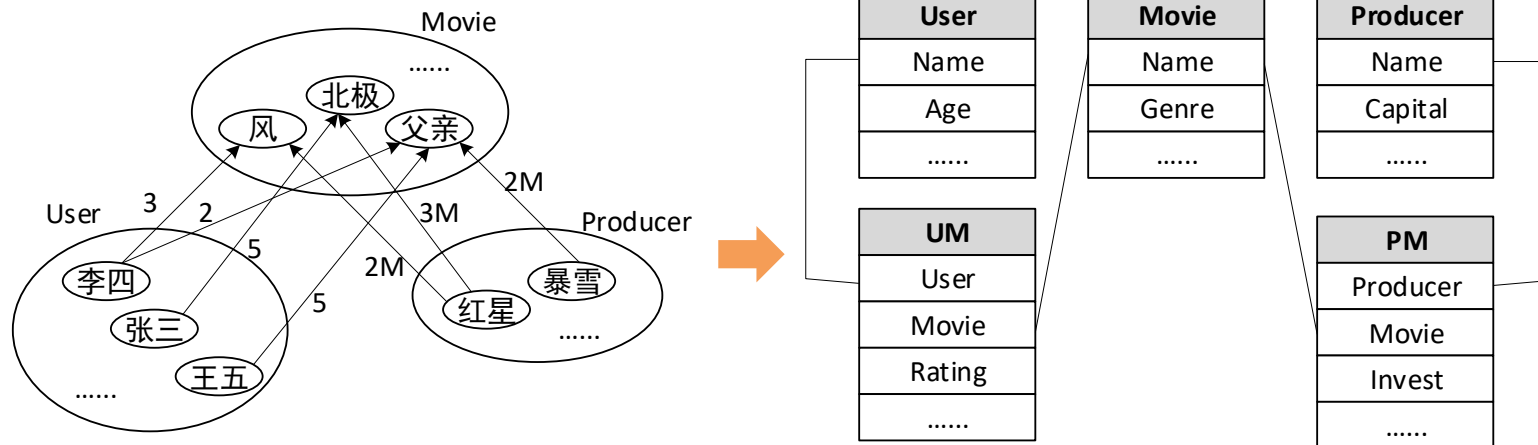
Mainly for graph mining algorithms

# Graph Query Output Pipeline (Compared)



# How to Get/Build a Graph

What does a graph look like in RDBMS?



How can we make the graph structure efficient?

System \ Algorithm	PageRank (5 iters)	Components
<b>Ligra</b>	6.69s	6.75s
<b>Neo4j</b>	131s	189s


Table 1: Performance Comparison (Twitter-2010)

We need CSR/CSC to be fast!

*CSR/C : Compressed Sparse Row/Column*

# CSR/C Explained

## Compressed sparse storage



Similar to Arrays in ClickHouse


## Vanilla edge list storage

0,1	0,2	0,3	0,4	1,2	3,2	3,4	4,1
-----	-----	-----	-----	-----	-----	-----	-----

Easy to store and modify (normal tables)

# From Tables to CSR/C

Define vertex tables to encode key columns



```
CREATE TABLE Producer (Name String key, Capital UInt64) ENGINE V;
CREATE TABLE Movie (Name String key, Genre String) ENGINE V;
CREATE TABLE User (Name String key, Age int) ENGINE V;
```

Key: HashMap -> [0, P-1]	Producer
	Name: Key
	Capital
Property: Columns	.....
Vertex Table: Producer	

红星: 0  
暴雪: 1  
E.g.

Key: HashMap -> [0, M-1]	Movie
	Name: Key
	Genre
Property: Columns	.....
Vertex Table: Movie	


风: 0  
北极: 1  
父亲: 2  
E.g.

Key: HashMap -> [0, N-1]	User
	Name: Key
	Age
Property: Columns	.....
Vertex Table: User	

张三: 0  
李四: 1  
王五: 2  
E.g.

# From Tables to CSR/C

## Define edge tables to store relations



```
CREATE TABLE Producer_Movie (Src VS(Producer), Dst VD(Movie), Invest UInt64) ENGINE E;  
CREATE TABLE User_Movie (Src VS(User), Dst VD(Movie), Rating int) ENGINE E;
```

src: _VS (Producer) dst: _VD (Movie) order by <_VS, _VD>  Property: Columns	Producer_Movie
	Producer: VS
	Movie: VD
	Invest
	.....
Edge Table: Producer_Movie	


红星-风: <0, 0>  
红星-北极: <0, 1>  
暴雪-父亲: <1, 1>  
E.g.

src: _VS (User) dst: _VD (Movie) order by <_VS, _VD>  Property: Columns	User_Movie
	User: VS
	Movie: VD
	Rating
	.....
Edge Table: User_Movie	

张三-北极: <0, 1>  
李四-风: <1, 0>  
李四-父亲: <1, 2>  
王五-父亲: <2, 2>  
E.g.

Storing hidden columns \_vs, \_vd and use 'MergeTree order by \_vs, \_vd' as the underlying engine

# From Tables to CSR/C



Define MetaSQL to specify how to build a graph

```
CREATE GRAPH User_Movie_Producer AS SELECT * FROM edgeGroup(User_Movie, Producer_Movie)
```


MetaSQL Defines the following workflow:

1. Adjusting IDs of vertices and edges
2. Aggregating edge properties
3. Grouping edges (also assigning Eid)

张三: 0	<table><tr><td>User:</td><td>Movie:</td><td>Producer:</td></tr><tr><td>vertex num=N</td><td>vertex num=M</td><td>vertex num=P</td></tr><tr><td>key=Name</td><td>key=Name</td><td>key=Name</td></tr><tr><td>Property Handle</td><td>Property Handle</td><td>Property Handle</td></tr></table>			User:	Movie:	Producer:	vertex num=N	vertex num=M	vertex num=P	key=Name	key=Name	key=Name	Property Handle	Property Handle	Property Handle	<table><tr><td>edge group</td></tr><tr><td>vs</td></tr><tr><td>vd</td></tr><tr><td>Rating</td></tr><tr><td>Invest</td></tr><tr><td>.....</td></tr></table>	edge group	vs	vd	Rating	Invest	.....	Eid
User:	Movie:	Producer:																					
vertex num=N	vertex num=M	vertex num=P																					
key=Name	key=Name	key=Name																					
Property Handle	Property Handle	Property Handle																					
edge group																							
vs																							
vd																							
Rating																							
Invest																							
.....																							
李四: 1				张三-北极: <0, 4>	1																		
王五: 2				李四-风: <1, 3>	2																		
风: 3				李四-父亲: <1, 5>	3																		
北极: 4				王五-父亲: <2, 5>	4																		
父亲: 5				红星-风: <6, 3>	5																		
红星: 6				红星-北极: <6, 4>	6																		
暴雪: 7				暴雪-父亲: <7, 5>	7																		
E.g.	Vertex Table Proxy			E.g.																			
	MetaSQL: User_Movie_Producer																						

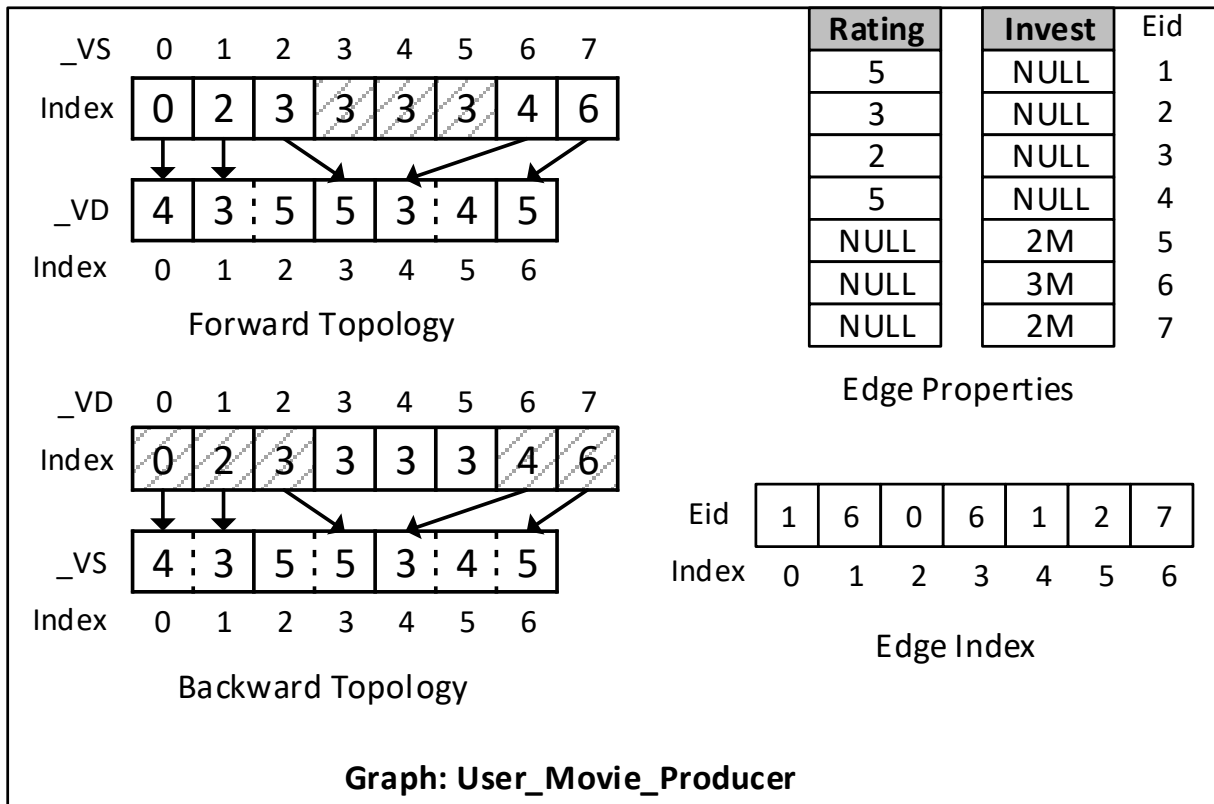


# From Tables to CSR/C



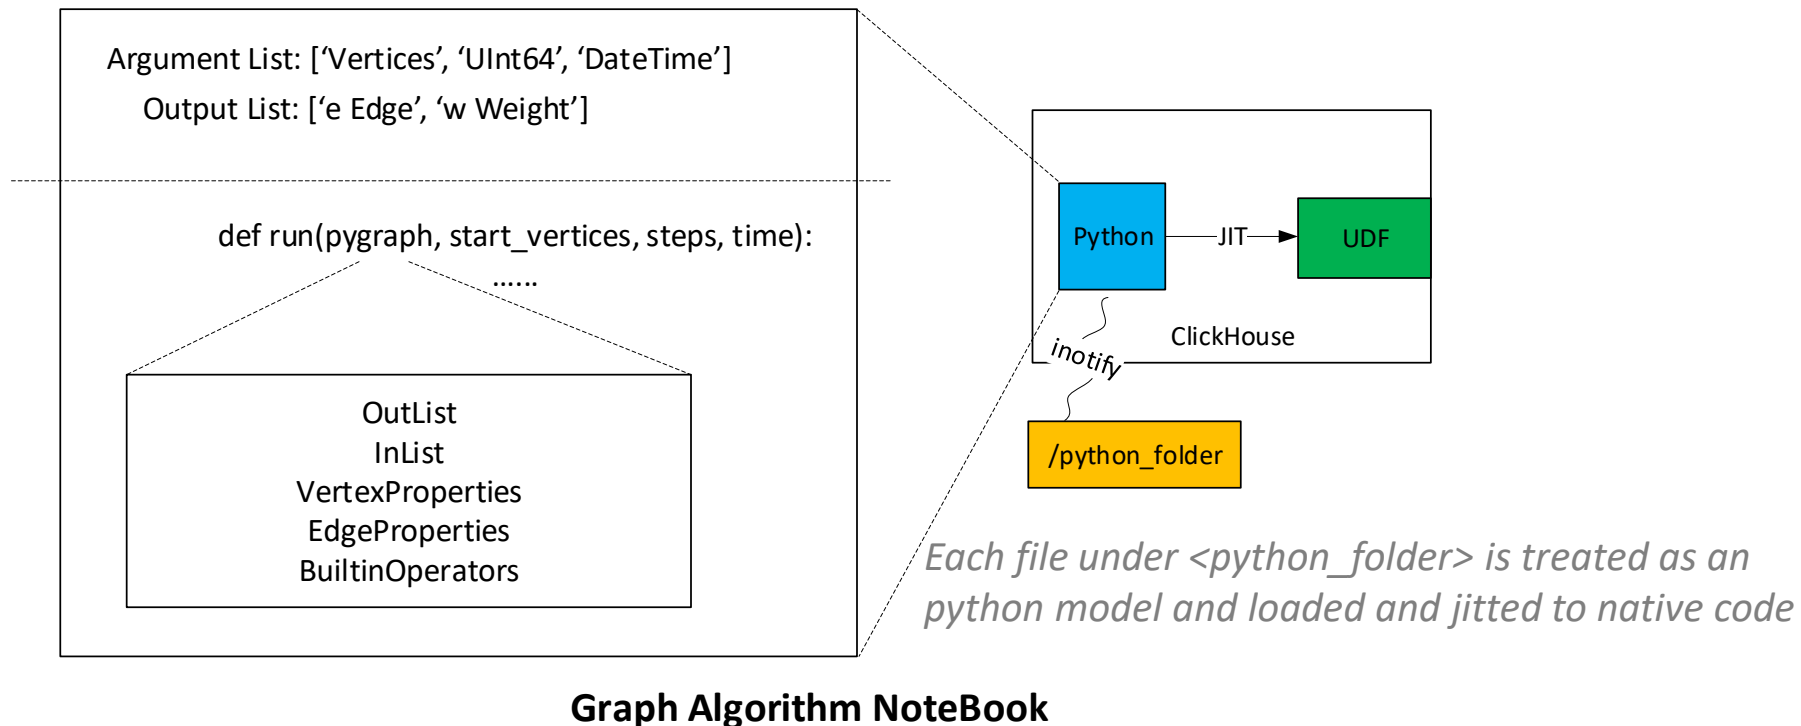
Execute MetaSQL to build a graph

*REFRESH User\_Movie\_Producer [FULL]*



# How to Write a Graph Algorithm

- Two ways of writing graph algorithms: python or c++
- **Python:** for ad-hoc queries, debugging and performance tuning

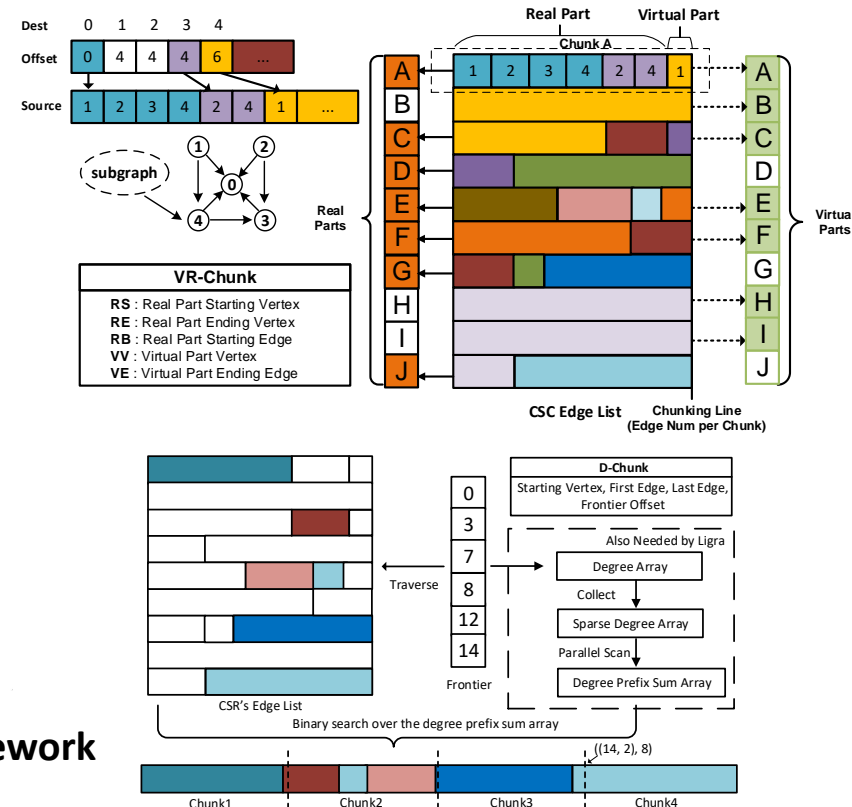


# How to Write a Graph Algorithm

- Two ways of writing graph algorithms: python or c++
- **C++**: for long use/well known algorithms, efficiency

```
void PageRankFunction::run(Graph & graph) {
    ... // initialization code
    Algorithm algo(
        [&](UInt32 s, UInt32 d) { // push
            atomicAdd(pr_new[d], pr[s]); },
        [&](UInt32* b, UInt32* e, UInt32 rd, UInt32 d) { // pull
            Float y = 0;
            while (b < e) y += pr[*b++];
            pr_new[d] = y; },
        [&](UInt32 rd, UInt32 d) { // pull reduce
            atomicAdd(pr_new[rd], pr_new[d]); }
    );
    while (!finish) {
        graph.run(algo);
        ... /* other related code */
    }
}
```

Listing 1. Page Rank Implementation



SilverChunk Framework

# Graph Query Helper

- vp, ep, pp for outputting related properties
- v for getting the internal id(s) of a vertex/vertices
- edge, vertex for retrieving edges/vertices of a graph respectively
- show graphs/algos for global inspections
- desc graph/algo for detail info
- graphInfo as a binary interface to access graph info

.....

# Graph Query Helper (E.g.)

DESCRIBE GRAPH wz

	name	value or type	comment
Edge properties:	region	String	vertex id range : 0 --- 15
	time	Date	
	w	Float64	
Vertex group:			
Vertex name:	default.node		
Key name:	id	UInt32	
Vertex properties:	id	UInt32	
	age	UInt32	
	name	String	
Meta info:	vertex_num	16	
	edge_num	16	
	symmetric	false	
	load	true	

# Graph Query Helper (E.g.)

DESCRIBE ALGORITHM pagerank

	name	value or type	comment
Arguments:		Vertices UInt64 Float64 Float64	vertices : start vertices (default all) iter : iterations (default 5), 0 for unlimited damping : damping value (default 0.85) epsilon : epsilon value (default 0.001)
Outputs:	v pagerank	Vertices Float32	

# Graph Query Helper (E.g.)

SHOW ALGOS

name	type	comment
bfs	python	
dijkstra	python	
vertexfilter	python	
simplebfs	builtin	
cc	builtin	
bellmanford	builtin	
coregroup	builtin	
linkcircle	builtin	
pagerank	builtin	
bc	builtin	
hop	builtin	
ssp	builtin	
shortcutcc	builtin	
personalrank	builtin	
triangle	builtin	
neighbor	builtin	
closec	builtin	
path	builtin	
commonneighbors	builtin	
radiiest	builtin	
	.....	

# Limitations

- Single Machine, Main Memory
- **But we have good reasons**
  - Graph size aren't that huge for us (around multiple billions of edges)
    - also commonly seen (referring Sahu, Siddhartha, et al PVLDB 17)
  - Main memory is big, very big (new Mac Pro supports 1.5TB RAM)
  - Very popular in graph computing community (Ligra, Galois, Grazelle, etc.)
  - Nearly impossible to have independent partitions for a graph
    - Commodity networks are too slow, InfiniBand is too expensive
    - Scalability! But at what COST? (McSherry HotOS15)

scalable system	cores	twitter	uk-2007-05
GraphChi [12]	2	3160s	6972s
Stratosphere [8]	16	2250s	-
X-Stream [21]	16	1488s	-
Spark [10]	128	857s	1759s
Giraph [10]	128	596s	1235s
GraphLab [10]	128	249s	833s
GraphX [10]	128	419s	462s
Single thread (SSD)	1	300s	651s
Single thread (RAM)	1	275s	-

20 PageRank iterations

scalable system	cores	twitter	uk-2007-05
Stratosphere [8]	16	950s	-
X-Stream [21]	16	1159s	-
Spark [10]	128	1784s	$\geq 8000s$
Giraph [10]	128	200s	$\geq 8000s$
GraphLab [10]	128	242s	714s
GraphX [10]	128	251s	800s
Single thread (SSD)	1	153s	417s

Label Propagation



# Future Works

- External Property Partitions
  - Sequential access to properties stored in external storages
- Partial/Incremental Graph Loading
  - Static knowledge database (huge graph) plus loadable domain data
- After the Processor branch Landed
  - Vectorized Graph Queries
  - Streaming and Time Series Processing

# Lessons Learned While Extending CH

- Understanding of code
  - semantic tools, utilization of debugger, reading unit tests
  - trial and error
- Isolation of code
  - new modules > new subclasses > Pimpl isolation > in-place patching
  - dlopen for optional features (e.g. python)
- Compilation (linking) is sloooooow
  - shared build might help
  - better to build a libclickhouse.so and write main functions per feature
- Contribute while extending

# Thank You!

Questions?